

```

# ===== INSTALL =====
!pip install -q transformers torch

# ===== IMPORTS =====
from transformers import pipeline, set_seed, GPT2LMHeadModel, GPT2Tokenizer
import torch

# ===== SET SEED =====
set_seed(42)

# ===== PIPELINE =====
generator = pipeline(
    'text-generation',
    model='gpt2',
    device=0 if torch.cuda.is_available() else -1
)

prompt = "Once upon a time in a distant galaxy,"

outputs = generator(
    prompt,
    max_new_tokens=100,
    num_return_sequences=3,
    temperature=0.8,
    top_p=0.95,
    do_sample=True,
    pad_token_id=50256
)

print("=== Text Generations ===")
for i, output in enumerate(outputs):
    print(f"\nGeneration {i+1}:\n{output['generated_text']}")
    print("-" * 80)

# ===== MANUAL =====
tokenizer = GPT2Tokenizer.from_pretrained('gpt2')
model = GPT2LMHeadModel.from_pretrained('gpt2')

tokenizer.pad_token = tokenizer.eos_token

prompt2 = "In a world where AI has become conscious,"

inputs = tokenizer(prompt2, return_tensors="pt").to(
    0 if torch.cuda.is_available() else "cpu"
)

outputs2 = model.generate(
    **inputs,
    max_new_tokens=120,
    temperature=0.7,
    top_p=0.9,
    do_sample=True,
    repetition_penalty=1.2
)

print("\n=== Manual Generation ===")
print(tokenizer.decode(outputs2[0], skip_special_tokens=True))

```

```

/usr/local/lib/python3.12/dist-packages/huggingface_hub/utils/_auth.py:93: UserWarning:
The secret `HF_TOKEN` does not exist in your Colab secrets.
To authenticate with the Hugging Face Hub, create a token in your settings tab (https://huggingface.co/settings/tokens), set it
You will be able to reuse this secret in all of your notebooks.
Please note that authentication is recommended but still optional to access public models or datasets.
warnings.warn(

Loading weights: 100% 148/148 [00:00<00:00, 265.83it/s, Materializing param=transformer.wte.weight]

GPT2LMHeadModel LOAD REPORT from: gpt2
Key | Status | |
-----+-----+
h.{0...11}.attn.bias | UNEXPECTED | |

Notes:
- UNEXPECTED :can be ignored when loading from different task/architecture; not ok if you expect identical arch.
Passing `generation_config` together with generation-related arguments=({'temperature', 'max_new_tokens', 'do_sample', 'pad_token_id'})
Both `max_new_tokens` (=100) and `max_length` (=50) seem to have been set. `max_new_tokens` will take precedence. Please refer to
=== Text Generations ===

Generation 1:
Once upon a time in a distant galaxy, a man in his thirties, in his early thirties, was brought to the knowledge that the universe
-----

Generation 2:
Once upon a time in a distant galaxy, there was nothing but the light of a light bulb and the sound of a radio. A light bulb was

It is a small light bulb, and a small light bulb is very small.

The smallest light bulb of any kind was the most valuable part of the universe.

The most valuable part of the universe was what the universe made of light.

The most valuable
-----

Generation 3:
Once upon a time in a distant galaxy, the galaxy will pass through a new stage where life will exist and its potential is limitless

This could have been a huge step forward for our understanding of how our universe works. It also provides some clues to how to
-----

Warning: You are sending unauthenticated requests to the HF Hub. Please set a HF_TOKEN to enable higher rate limits and faster downloads
WARNING:huggingface_hub.utils._http:Warning: You are sending unauthenticated requests to the HF Hub. Please set a HF_TOKEN to enable

Loading weights: 100% 148/148 [00:00<00:00, 453.86it/s, Materializing param=transformer.wte.weight]

GPT2LMHeadModel LOAD REPORT from: gpt2
Key | Status | |
-----+-----+
h.{0...11}.attn.bias | UNEXPECTED | |

Notes:
- UNEXPECTED :can be ignored when loading from different task/architecture; not ok if you expect identical arch.
Setting `pad_token_id` to `eos_token_id`:50256 for open-end generation.

=== Manual Generation ===
In a world where AI has become conscious, we can only imagine how much of an impact it will have.
The more I thought about this subject in my mind the clearer and clearer what happened to me as you began your research into artificial

```